

In this project, I will build a machine learning model which will predict house price.



Firstly, I will load the data and clean the data and check the data set.

### ▼ Default title text

```
1 # @title Default title text
2 import numpy as np
3 import pandas as pd
4 import matplotlib.pyplot as plt
5 import seaborn as sns
6
7 # Load dataset from CSV (replace with your actual file path)
8 df = pd.read_csv("california_housing.csv")
9
10 # Check the first few rows
11 df.head()
12
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value	ocean_proximity
0	-122.23	37.88	41.0	880.0	129.0	322.0	126.0	8.3252	452600.0	NEAR BAY
1	-122.22	37.86	21.0	7099.0	1106.0	2401.0	1138.0	8.3014	358500.0	NEAR BAY
2	-122.24	37.85	52.0	1467.0	190.0	496.0	177.0	7.2574	352100.0	NEAR BAY
3	-122.25	37.85	52.0	1274.0	235.0	558.0	219.0	5.6431	341300.0	NEAR BAY
4	-122.25	37.85	52.0	1627.0	280.0	565.0	259.0	3.8462	342200.0	NEAR BAY

Next steps: [View recommended plots](#) [New interactive sheet](#)

```
1 df.info()
2 df['ocean_proximity'].unique()
```

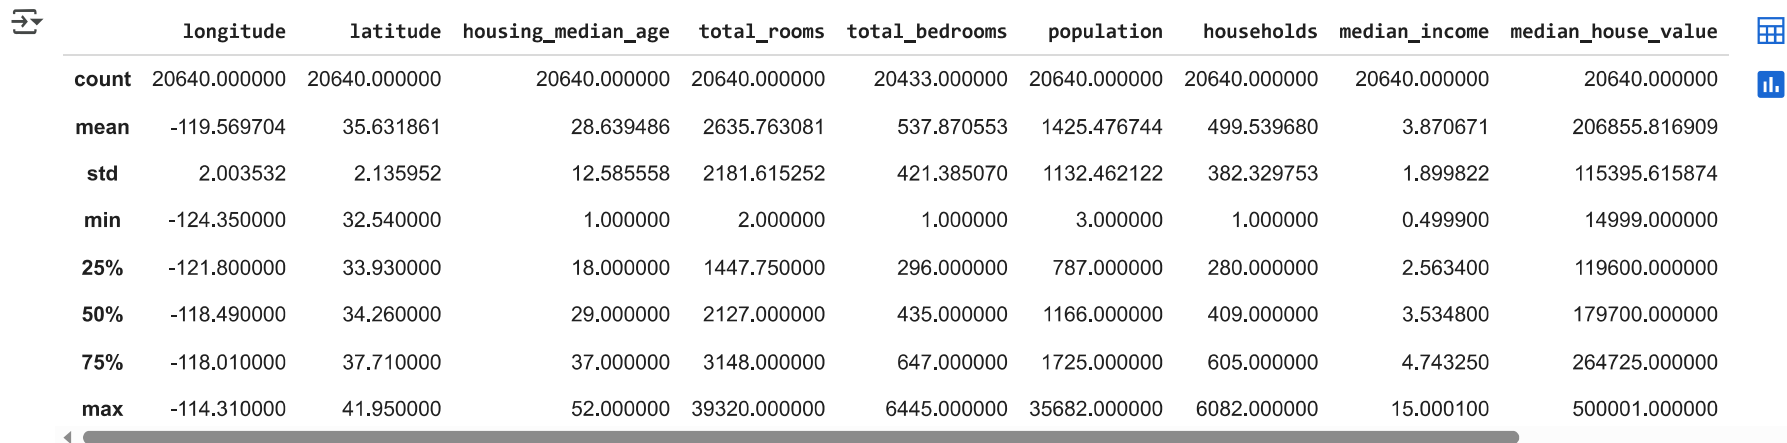
```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
```

```

#   Column          Non-Null Count  Dtype
---  -
0   longitude        20640 non-null  float64
1   latitude          20640 non-null  float64
2   housing_median_age 20640 non-null  float64
3   total_rooms        20640 non-null  float64
4   total_bedrooms     20433 non-null  float64
5   population         20640 non-null  float64
6   households         20640 non-null  float64
7   median_income      20640 non-null  float64
8   median_house_value 20640 non-null  float64
9   ocean_proximity    20640 non-null  object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB
array(['NEAR BAY', '<1H OCEAN', 'INLAND', 'NEAR OCEAN', 'ISLAND'],
      dtype=object)

```

```
1 df.describe()
```



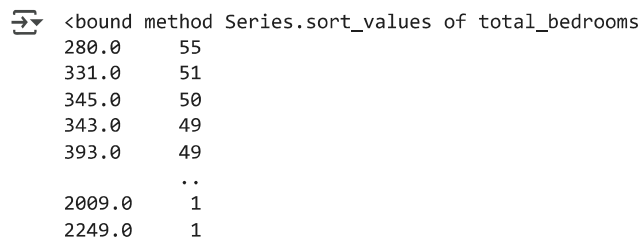
The table displays the statistical summary of the dataset. It includes counts, mean, standard deviation, minimum, and maximum values for each column, as well as the 25th, 50th (median), and 75th percentiles.

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value
<b>count</b>	20640.000000	20640.000000	20640.000000	20640.000000	20433.000000	20640.000000	20640.000000	20640.000000	20640.000000
<b>mean</b>	-119.569704	35.631861	28.639486	2635.763081	537.870553	1425.476744	499.539680	3.870671	206855.816909
<b>std</b>	2.003532	2.135952	12.585558	2181.615252	421.385070	1132.462122	382.329753	1.899822	115395.615874
<b>min</b>	-124.350000	32.540000	1.000000	2.000000	1.000000	3.000000	1.000000	0.499900	14999.000000
<b>25%</b>	-121.800000	33.930000	18.000000	1447.750000	296.000000	787.000000	280.000000	2.563400	119600.000000
<b>50%</b>	-118.490000	34.260000	29.000000	2127.000000	435.000000	1166.000000	409.000000	3.534800	179700.000000
<b>75%</b>	-118.010000	37.710000	37.000000	3148.000000	647.000000	1725.000000	605.000000	4.743250	264725.000000
<b>max</b>	-114.310000	41.950000	52.000000	39320.000000	6445.000000	35682.000000	6082.000000	15.000100	500001.000000

```

1 df.isnull().sum()
2 # Get the unique values and their frequency in the 'total_bedrooms' column
3 bedroom_counts = df['total_bedrooms'].value_counts().sort_values
4
5 # Display the results
6 print(bedroom_counts)

```



The output shows the frequency of each total bedroom count, sorted in ascending order. The counts range from 55 for 280 bedrooms to 1 for 2249 bedrooms.


total_bedrooms	count
280.0	55
331.0	51
345.0	50
343.0	49
393.0	49
..	..
2009.0	1
2249.0	1

```
3114.0    1
1625.0    1
1742.0    1
Name: count, Length: 1923, dtype: int64>
```

As I assume that a house must have at least 1 bedroom and therefore, I substitute 1 at the places of null values in bedroom column. And as you can see in the box plot, there are no outliers.

```
1 # Fill missing values in 'total_bedrooms' column with 1
2 df['total_bedrooms'] = df['total_bedrooms'].fillna(1)
3
```

```
1 # Visualize data distribution to check for outliers
2 sns.boxplot(data=df)
3 plt.xticks(rotation=90)
4 plt.show()
```



Preparing for XGBoost Machine Learning.

- Binary columns are created for 5 types of variables in ocean\_proximity column.
- We split the data into 80% training and 20% test groups. Here I use Standard Scaler to standardize the features in dataset by transforming data to have a mean of 0 and standard deviation of 1

```

1 from sklearn.model_selection import train_test_split
2 from sklearn.preprocessing import StandardScaler
3
4 # Split data into features (X) and target (y)
5 X = df.drop('median_house_value', axis=1)
6 y = df['median_house_value']
7
8 # Identify categorical columns
9 categorical_columns = X.select_dtypes(include=['object']).columns
10
11 # Perform One-Hot Encoding for categorical columns
12 X_encoded = pd.get_dummies(X, columns=categorical_columns)
13
14
15 # Split the data into training and testing sets
16 X_train, X_test, y_train, y_test = train_test_split(X_encoded, y, test_size=0.2, random_state=42)
17
18 # Standardize the features
19 scaler = StandardScaler()
20 X_train = scaler.fit_transform(X_train)
21 X_test = scaler.transform(X_test)
22
23 X_encoded.head()

```



	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	ocean_proximity_<1H OCEAN	ocean_proximity_INLAND	ocean_proximity_I
0	-122.23	37.88	41.0	880.0	129.0	322.0	126.0	8.3252	False	False	
1	-122.22	37.86	21.0	7099.0	1106.0	2401.0	1138.0	8.3014	False	False	
2	-122.24	37.85	52.0	1467.0	190.0	496.0	177.0	7.2574	False	False	
3	-122.25	37.85	52.0	1274.0	235.0	558.0	219.0	5.6431	False	False	
4	-122.25	37.85	52.0	1627.0	280.0	565.0	259.0	3.8462	False	False	

Next steps:

[View recommended plots](#)

[New interactive sheet](#)

Now, I am building a machine learning model with XG Boost. In this model, we are doing regression with 'reg:squarederror' as the loss function and the number of boosting rounds or tree is 1000 and learning 0.05 and the maximum depth of tree is 6 as default hyperparameter values.

```

1 import xgboost as xgb
2 from sklearn.metrics import mean_squared_error, r2_score
3
4 # Initialize the XGBoost regressor
5 model = xgb.XGBRegressor(objective='reg:squarederror', n_estimators=1000, learning_rate=0.05, max_depth=6)
6
7 # Train the model
8 model.fit(X_train, y_train)
9
10 # Make predictions
11 y_pred = model.predict(X_test)
12
13 # Evaluate the model
14 mse = mean_squared_error(y_test, y_pred)
15 r2 = r2_score(y_test, y_pred)
16
17 print(f'Mean Squared Error: {mse}')
18 print(f'R-squared: {r2}')
19

```

➡ Mean Squared Error: 2687655514.744204  
R-squared: 0.7948994886808447

From the result above you can see R-squared value is 0.79 and I use cross validation to check the result furthermore and you can see the overall R squared value is 0.83 which means 83% of the model variation is explained by my model features and 18% remained unexplained and you can say the model is performing well

```


1 from sklearn.model_selection import cross_val_score
2 scores = cross_val_score(model, X_train, y_train, cv=5, scoring="r2")
3 print(f"Cross-Validation R² Scores: {scores}")
4 print(f"Average R² Score: {scores.mean()}")

```

➡ Cross-Validation R² Scores: [0.84486807 0.83278214 0.84203853 0.83943947 0.84017939]  
Average R² Score: 0.8398615215417877

Below is the residual plot, which helps evaluate the model's performance. The residuals (differences between actual and predicted values) are scattered randomly around the red reference line. Since there is no clear pattern in the residual distribution, we can conclude that the model is performing well, indicating that the errors are randomly distributed.

```
1 # Calculate the residuals (differences between actual and predicted values)
2 residuals = y_test - y_pred
3
4 # Plot residuals
5 plt.figure(figsize=(10, 6))
6 sns.scatterplot(x=y_test, y=residuals, color='purple')
7
8 plt.axhline(0, color='red', linestyle='--') # Ideal line at y = 0
9 plt.title('Residuals vs Actual House Prices')
10 plt.xlabel('Actual House Price')
11 plt.ylabel('Residuals (Actual - Predicted)')
12 plt.show()
```



After training the model, I am going to optimize tuning hyper parameters like `n_estimators`, `learning_rate`, `max_depth` using tools like `GridSearchCV`

```
1 from sklearn.model_selection import GridSearchCV
2
```

```

3 # Define hyperparameters to tune
4 param_grid = {
5     'n_estimators': [1000, 1200, 1500],
6     'learning_rate': [0.1, 0.3, 0.4, 0.5],
7     'max_depth': [3, 6, 10]
8 }
9
10 # Perform Grid Search with cross-validation
11 grid_search = GridSearchCV(xgb.XGBRegressor(objective='reg:squarederror'), param_grid, cv=3, scoring='neg_mean_squared_error')
12 grid_search.fit(X_train, y_train)
13
14 # Best parameters from grid search
15 print("Best Parameters: ", grid_search.best_params_)
16

```

Best Parameters: {'learning\_rate': 0.1, 'max\_depth': 6, 'n\_estimators': 1200}

So, I am training a new model with best parameters. {'learning\_rate': 0.1, 'max\_depth': 6, 'n\_estimators': 1200}

```


1 import xgboost as xgb
2 from sklearn.metrics import mean_squared_error, r2_score
3
4 # Initialize the XGBoost regressor
5 model = xgb.XGBRegressor(objective='reg:squarederror', n_estimators=1200, learning_rate=0.1, max_depth=6)
6
7 # Train the model
8 model.fit(X_train, y_train)
9
10 # Make predictions
11 y_pred = model.predict(X_test)
12
13 # Evaluate the model
14 mse = mean_squared_error(y_test, y_pred)
15 r2 = r2_score(y_test, y_pred)
16
17 print(f'Mean Squared Error: {mse}')
18 print(f'R-squared: {r2}')
19

```

Mean Squared Error: 3022138276.9204564  
R-squared: 0.7693744222527084

Now, I am visualizing the actual and predicted to see how well they match. You can see the dots are even spread throughout the red line and the model is well performing.

```
1 # Create a DataFrame to hold the actual vs predicted values
2 comparison_df = pd.DataFrame({'Actual': y_test, 'Predicted': y_pred})
3
4 # Plot the actual vs predicted values
5 plt.figure(figsize=(10, 6))
6 sns.scatterplot(data=comparison_df, x='Actual', y='Predicted', color='blue')
7
8 # Add a line of perfect prediction (y = x)
9 plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='red', linestyle='--')
10
11 plt.title('Actual vs Predicted House Prices')
12 plt.xlabel('Actual House Price')
13 plt.ylabel('Predicted House Price')
14 plt.show()
15
16
```






Instead of improving the model's performance, the R-squared value decreased from 0.79 to 0.76, indicating a slight decline in accuracy. However, the residual plot shows that the residuals are more tightly clustered around the red reference line. This suggests that while the model's variance has reduced, it may have lost some predictive power

```
1 # Calculate the residuals (differences between actual and predicted values)
2 residuals = y_test - y_pred
3
4 # Plot residuals
5 plt.figure(figsize=(10, 6))
6 sns.scatterplot(x=y_test, y=residuals, color='purple')
7
8 plt.axhline(0, color='red', linestyle='--') # Ideal line at y = 0
9 plt.title('Residuals vs Actual House Prices')
10 plt.xlabel('Actual House Price')
11 plt.ylabel('Residuals (Actual - Predicted)')
12 plt.show()
13
```



As shown in the results, the cross-validation R-squared mean score has also decreased by approximately 0.0018 compared to the model with default parameters. Since the tuned model does not provide significant improvement, I will proceed with the default parameters

```
1 from sklearn.model_selection import cross_val_score
2 scores = cross_val_score(model, X_train, y_train, cv=5, scoring="r2")
3 print(f"Cross-Validation R² Scores: {scores}")
4 print(f"Average R² Score: {scores.mean()}")
```

 Cross-Validation R² Scores: [0.84521146 0.83238389 0.84013603 0.83674848 0.83552197]  
Average R² Score: 0.838000367965028